

Assignment 3 - Multi-class Classification and Neural Network

FULL MARKS = 10

In this assignment, you are going to implement your own neural network to do multi-class classification. We use one-vs-all strategy here by training multiple binary classifiers (one for each class).

Please notice that you can only use numpy and scipy.optimize.minimize. No library versions of other method are allowed. . Follow the instructions, you will need to fill the blanks to make it functional. The process is similar to the previous assignment.

```
In [1]: from sklearn import datasets
        from scipy.optimize import minimize
        import numpy as np
```

```
In [2]: def load_dataset():
        iris = datasets.load_iris()
        X = iris.data
        y = iris.target

        return X, y
```

```
In [3]: def train_test_split(X, y):
        idx = np.arange(len(X))
        train_size = int(len(X) * 2/3)
        np.random.shuffle(idx)
        X = X[idx]
        y = y[idx]
        X_train, X_test = X[:train_size], X[train_size:]
        y_train, y_test = y[:train_size], y[train_size:]

        return X_train, X_test, y_train, y_test
```

```
In [4]: def init_weights(num_in, num_out):
        """
        :param num_in: the number of input units in the weight array
        :param num_out: the number of output units in the weight array
        """

        # Note that 'W' contains both weights and bias, you can consider W[0, :] as bias
        W = np.zeros((1 + num_in, num_out))

        #####
        # Full Mark: 1
        # TODO:
        # Implement Xavier/Glorot uniform initialization
        #
        # Hint: you can find the reference here:
        # https://www.tensorflow.org/api_docs/python/tf/keras/initializers/GlorotUniform
        #####

        x_uniform = np.sqrt(6.0 / (num_in + num_out))
        W = np.random.uniform(-x_uniform, x_uniform, W.shape)

        #####
        #                               END OF YOUR CODE
        #####

        return W
```

```
In [5]: def sigmoid(x):
        """
        :param x: input
        """

        #####
        # Full Mark: 0.5
        # TODO:
        # Implement sigmoid function:
        #                               sigmoid(x) = 1/(1+e^(-x))
        #####

        res = 1 / ( 1 + np.exp(-x))

        #####
        #                               END OF YOUR CODE
        #####

        return res
```

```
In [6]: def tanh(x):
        """
```

```

:param x: input
'''

#####
# Full Mark: 0.5
# TODO:
# Implement tanh function:
#  $\tanh(x) = (e^x - e^{-x}) / (e^x + e^{-x})$ 
#####

res = np.tanh(x)

#####
#
# END OF YOUR CODE
#####

return res

```

```

In [7]: def sigmoid_gradient(x):
'''
:param x: input
'''

#####
# Full Mark: 1
# TODO:
# Computes the gradient of the sigmoid function evaluated at x.
#
#####

grad = sigmoid(x)*(1-sigmoid(x))

#####
#
# END OF YOUR CODE
#####

return grad

```

```

In [8]: def tanh_gradient(x):
'''
:param x: input
'''

#####
# Full Mark: 1
# TODO:
# Computes the gradient of the tanh function evaluated at x.
#
#####

grad = (1 - tanh(x)**2)

#####
#
# END OF YOUR CODE
#####

return grad

```

```

In [9]: def forward(W, X):
'''
:param W: weights (including biases) of the neural network. It is a list containing both W_hidden and W_output
:param X: input. Already added one additional column of all "1"s.
'''

# Shape of W_hidden: [num_feature+1, num_hidden]
# Shape of W_output: [num_hidden+1, num_output]
W_hidden, W_output = W

#####
# Full Mark: 1
# TODO:
# Implement the forward pass. You need to compute four values:
# z_hidden, a_hidden, z_output, a_output
#
# Note that our neural network consists of three layers:
# Input -> Hidden -> Output
# The activation function in hidden layer is 'tanh'
# The activation function in output layer is 'sigmoid'
#####

# convert to numpy array
W_hidden = np.array(W_hidden)
W_output = np.array(W_output)

# Forward pass through hidden layer
z_hidden = np.dot(X, W_hidden)
a_hidden = tanh(z_hidden)

```

```

# Add bias column
a_hidden = np.concatenate([np.ones(( X.shape[0], 1)), a_hidden], axis=1)

# Forward pass through Output layer
z_output = np.dot(a_hidden,W_output)
a_output = sigmoid(z_output)

#####
#                               END OF YOUR CODE                               #
#####

# z_hidden is the raw output of hidden layer, a_hidden is the result after performing activation on z_hidden
# z_output is the raw output of output layer, a_output is the result after performing activation on z_output
return {'z_hidden': z_hidden, 'a_hidden': a_hidden,
        'z_output': z_output, 'a_output': a_output}

```

```

In [10]: def loss_funtion(W, X, y, num_feature, num_hidden, num_output, L2_lambda):
'''
:param W: a 1D array containing all weights and biases.
:param X: input
:param y: labels of X
:param num_feature: number of features in X
:param num_hidden: number of hidden units
:param num_output: number of output units
:param L2_lambda: the coefficient of regularization term
'''

m = y.size

# Reshape W back into the parameters W_hidden and W_output
W_hidden = np.reshape(W[:num_hidden * (num_feature + 1)],
                      ((num_feature + 1), num_hidden))

W_output = np.reshape(W[(num_hidden * (num_feature + 1)):],
                      ((num_hidden + 1), num_output))

# Initialize grads
W_hidden_grad = np.zeros(W_hidden.shape)
W_output_grad = np.zeros(W_output.shape)

# Add one column of "1"s to X
X_input = np.concatenate([np.ones((m, 1)), X], axis=1)

#####
# Full Mark: 3
# TODO:
# 1. Transform y to one-hot encoding. Implement binary cross-entropy loss function
# (Hint: Use the forward function to get necessary outputs from the model)
#
# 2. Add L2 regularization to all weights in loss
# (Note that we will not add regularization to bias)
#
# 3. Compute the gradient for W_hidden and W_output (including both weights and biases)
# (Hint: use chain rule, and the sigmoid gradient, tanh gradient you have
# implemented above. Don't forget to add the gradient of regularization term)
#####

# One hot encoding
shape = (y.size, y.max()+1)
one_hot = np.zeros(shape)
rows = np.arange(y.size)
one_hot[rows, y] = 1

# forward pass
f_pass = forward((W_hidden,W_output),X_input)
a_hidden = f_pass['a_hidden']
a_output = f_pass['a_output']

# cross-entropy loss
logprobs = np.multiply(one_hot, np.log(f_pass['a_output'])) + np.multiply((1 - one_hot), np.log(1 - f_pass['a_output']))
L = (-1.0/m) * np.sum(logprobs)

# L2 regularize
sum_w = np.sum(np.square(W_hidden)) + np.sum(np.square(W_output))
L = L + (sum_w * (L2_lambda/(2*m)))
L = np.squeeze(L)

##
### Chain Rule
##

# Gradient at output layer
dZ2 = sigmoid_gradient(a_output)
dW2 = np.dot(a_hidden[:,1:].T,dZ2)/m + ((L2_lambda / m) * W_output[1,:])
dB2 = np.sum(dZ2, axis=0)/m

# Gradient at hidden layer

```

```

dZ1 = np.multiply(dZ2.dot(W_output[1:,:].T), tanh_gradient(a_hidden[:,1:]))
dW1 = (np.dot(X.T, dZ1))/m + ((L2_lambda / m) * W_hidden[1:,:])
dB1 = np.sum(dZ1, axis=0)/m

# Reshape Bias Gradient matrix
dB1 = dB1.reshape(1,10)
dB2 = dB2.reshape(1,3)

# Final Gradients of weights
W_hidden_grad = np.concatenate((dB1, dW1), axis = 0)
W_output_grad = np.concatenate((dB2, dW2), axis = 0)

#####
#                               END OF YOUR CODE                               #
#####

grads = np.concatenate([W_hidden_grad.ravel(), W_output_grad.ravel()])

return L, grads

```

```

In [11]: def optimize(initial_W, X, y, num_epoch, num_feature, num_hidden, num_output, L2_lambda):
'''
:param initial_W: initial weights as a 1D array.
:param X: input
:param y: labels of X
:param num_epoch: number of iterations
:param num_feature: number of features in X
:param num_hidden: number of hidden units
:param num_output: number of output units
:param L2_lambda: the coefficient of regularization term
'''

options = {'maxiter': num_epoch}

#####
# Full Mark: 1
# TODO:
# Optimize weights
# (Hint: use scipy.optimize.minimize to automatically do the iteration.)
# ref: https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.minimize.html)
# For some optimizers, you need to set 'jac' as True.
# You may need to use lambda to create a function with one parameter to wrap the
# loss_funtion you have implemented above.
#
# Note that scipy.optimize.minimize only accepts a 1D weight array as initial weights,
# and the output optimized weights will also be a 1D array.
# That is why we unroll the initial weights into a single long vector.
#####

# initialize W_final
W_final = np.copy(initial_W)

# Function to run at each epoch
def fun(W):
    global W_final
    L, grad = loss_funtion(W, X, y, num_feature, num_hidden, num_output, L2_lambda)
    W_final = W - grad
    return L

# Optimize (Run Epochs)
minimize(fun, initial_W, options=options)

#####
#                               END OF YOUR CODE                               #
#####

# Obtain W_hidden and W_output back from W_final
W_hidden = np.reshape(W_final[:num_hidden * (num_feature + 1)],
                      ((num_feature + 1), num_hidden))
W_output = np.reshape(W_final[(num_hidden * (num_feature + 1)):],
                      ((num_hidden + 1), num_output))

return [W_hidden, W_output]

```

```

In [12]: def predict(X_test, y_test, W):
'''
:param X_test: input
:param y_test: labels of X_test
:param W: a list containing two weights W_hidden and W_output.
'''

test_input = np.concatenate([np.ones((y_test.size, 1)), X_test], axis=1)

#####
# Full Mark: 1
# TODO:
# Predict on test set and compute the accuracy.
# (Hint: use forward function to get predicted output)
#

```

```

#
#####

# Forward pass
cache = forward(W, test_input)

# One-hot-encode actual values
shape = (y_test.size, y_test.max()+1)
one_hot = np.zeros(shape)
rows = np.arange(y_test.size)
one_hot[rows, y_test] = 1

# Encode predicted outputs
y_hat = np.zeros(shape)
for i, r in enumerate(cache['a_output']):
    max_val = np.amax(r)
    for j, v in enumerate(r):
        if max_val == v:
            y_hat[i,j] = 1

# Find Accuracy
acc = 0
for i, _ in enumerate(one_hot):
    if np.array_equal(one_hot[i], y_hat[i]):
        acc += 1

acc = acc / y_test.size

#####
#                               END OF YOUR CODE                               #
#####

return acc

```

```

In [13]: # Do not modify this part #
# Define parameters
NUM_FEATURE = 4
NUM_HIDDEN_UNIT = 10
NUM_OUTPUT_UNIT = 3
NUM_EPOCH = 100
L2_lambda = 1

# Load data
X, y = load_dataset()
X_train, X_test, y_train, y_test = train_test_split(X, y)

# Initialize weights
initial_W_hidden = init_weights(NUM_FEATURE, NUM_HIDDEN_UNIT)
initial_W_output = init_weights(NUM_HIDDEN_UNIT, NUM_OUTPUT_UNIT)
initial_W = np.concatenate([initial_W_hidden.ravel(), initial_W_output.ravel()], axis=0)

# Neural network learning
W = optimize(initial_W, X_train, y_train, NUM_EPOCH, NUM_FEATURE, NUM_HIDDEN_UNIT, NUM_OUTPUT_UNIT, L2_lambda)

# Predict
acc = predict(X_test, y_test, W)
print("Test accuracy:", acc)

```

Test accuracy: 0.74